



Multi-model WBE-Agent evaluation: performance and experimental workload

Each point or bar corresponds to one full agent run over the labeled anomaly benchmark. Workload is quantified from saved investigation outputs, tokens, and tool usage.



This figure emphasizes that the model-swap experiment is not a single benchmark pass. It is a repeated, tool-using agent evaluation over the same 199-event benchmark.